

Kunj P. Shah

+1 (628) 529-6990 | kunjcr2@gmail.com | [linkedin.com/in/kunjcr2](https://www.linkedin.com/in/kunjcr2) | github.com/kunjcr2 | kunjcr2.github.io
San Francisco, CA

F-1 international student eligible for CPT and OPT. Eligible for H-1B change of status; the \$100,000 H-1B fee does not apply.

Education & Certificates

- B.S. in Computer Science from San Jose State University, GPA: 3.88/4.00 Graduating May 2027

Skills

Python, SQL, Java, PyTorch, TensorFlow, Scikit-learn, Azure ML, Weights & Biases, Pandas, NumPy, LangChain, Transformers, Large Language Models, Fine Tuning, vLLM, FastAPI, Docker, PostgreSQL, CI/CD, Agile

Work History/Projects/Internships

Deep Learning Engineer Intern, A10 Networking, CA

May 2026 to Aug 2026

- Developed a multimodal machine learning guardrail system for VLM-powered applications, collaborating cross-functionally within a 3-person AI Security team to deliver production-ready model safety tooling.
- Trained and fine-tuned a compact language model using Python, PyTorch, and SFT with chain-of-thought rationales, achieving sub-200ms inference on Hopper GPUs via vLLM to outperform leading public guardrail benchmarks.
- Built a multi-stage data annotation pipeline using Python and automated quality filtering, producing a curated training corpus with inter-annotator agreement metrics and versioning code with Git.
- Designed a systematic evaluation framework with Weights & Biases experiment tracking, communicating technical diagnostics across safety policy categories to non-technical stakeholders.

AI Engineer Intern, Routes Technologies, Remote

Oct 2025 to Jan 2026

- Architected AI and machine learning pipelines using Python, SQL, and Azure Cloud, deploying models in production across 5 managed online endpoints with CI/CD and Weights & Biases observability.
- Built an LLM-powered NL-to-SQL pipeline with few-shot prompting, input normalization (470+ synonym mappings), and parameterized SQL sanitization, serving Generative AI features to end users.
- Developed an SVD collaborative filtering recommendation system on Azure Cloud using Python, scikit-learn, and Pandas, evaluated with NDCG@5 and Precision/Recall@K metrics via a nightly ETL job.
- Engineered multi-source data ingestion using Python, Scrapy, and GPT-4o-mini classification, taking direction from cross-functional stakeholders to align data pipelines with product requirements.

AI/ML Engineering Intern, Dreamable Inc., CA

May 2025 to Aug 2025

- Trained, tuned, and deployed machine learning models in a production Cloud environment using Python, PyTorch, TensorFlow, and Hugging Face, working collaboratively with a cross-functional team.
- Fine-tuned the Qwen-2.5 7B LLM on GCP Cloud using LoRA to lower training cost while delivering a Q&A model with comparable accuracy within budget, tracking experiments with Weights & Biases.
- Curated NLP datasets using Python, Pandas, NumPy, and the Hugging Face Datasets library, cleaning and preprocessing training data to improve data quality for model fine-tuning.

- Developed a Generative AI outreach agent using LangChain and the OpenAI API to automate messaging workflows, cutting manual outreach time and increasing response rates.

StableLM-2-1.6B End-to-End Post-Pretraining, Personal Project, San Francisco Jan 2026 to Feb 2026

- Built an end-to-end machine learning pipeline for training, tuning, and deploying an LLM in production using Python, PyTorch, and Transformers, fine-tuning StableLM 1.6B on 140K conversations via LoRA.
- Implemented GRPO alignment using TRL and a DeBERTa-v3 reward model to safety-align the LLM on 4K samples, publishing final model weights to HuggingFace for public inference.
- Developed a production inference API using Python, FastAPI, and vLLM, containerized with Docker and CI/CD for GPU-enabled Cloud deployment.