

Kunj P. Shah

+1 (628) 529-6990 | kunjcr2@gmail.com | linkedin.com/in/kunjcr2 | github.com/kunjcr2 | kunjcr2.github.io | San Francisco, CA

F-1 international student eligible for CPT and OPT. Eligible for H-1B change of status; the \$100,000 H-1B fee does not apply.

Education

- B.S. in Computer Science from San Jose State University, GPA: 3.88/4.00 (Dean's List) Graduating May 2027

Work Experience

Deep Learning Engineer Intern, A10 Networking, CA

May 2026 to Aug 2026

- Built and deployed a Multimodal Machine Learning guardrail system using Python and PyTorch across text, Image+text, and multi-turn modalities, serving production inference via In-house Hopper GPUs infrastructure and collaborating cross-functionally with a 5-person AI Security team.
- Trained and fine-tuned two separate LLM guardrail models (text-only and Multimodal) using SFT, and chain-of-thought rationales, then performed SLERP model merging to achieve unified Generative AI safety coverage without retraining on combined data.
- Deployed the merged model to production via vLLM on Hopper GPUs with CI/CD, achieving sub-200ms inference latency and outperforming leading public guardrail benchmarks across all safety policy categories.
- Built a multi-stage data annotation and feature engineering pipeline using Python, Pandas, Huggingface and automated quality filtering, producing a curated training corpus with inter-annotator agreement metrics, versioned with Git.
- Designed a systematic evaluation framework using Weights & Biases, communicating model performance diagnostics and technical results to non-technical stakeholders to drive safety policy decisions.

AI Engineer Intern, Routes Technologies, Remote

Oct 2025 to Jan 2026

- Architected AI and Machine Learning pipelines using Python, SQL, and Azure Cloud, deploying models in production across 5 managed online endpoints with CI/CD and Weights & Biases observability.
- Built an LLM-powered NL-to-SQL pipeline with few-shot prompting, input normalization (470+ synonym mappings), and parameterized SQL sanitization, serving Generative AI features to end users.
- Developed an SVD collaborative filtering recommendation system on Azure Cloud using Python, scikit-learn, and Pandas, evaluated with NDCG@5 and Precision/Recall@K metrics via a nightly ETL job.
- Engineered multi-source data ingestion using Python, Scrappy, and GPT-4o-mini classification, taking direction from cross-functional stakeholders to align data pipelines with product requirements.
- Built an ingredient normalization pipeline using Python, sentence-transformer embeddings (all-MiniLM-L6-v2), and a tiered matching strategy (exact match → 5,400+ synonym table → cosine similarity), deployed as chained Azure ML pipelines with automated health monitoring and a human-in-the-loop review queue, achieving 99%+ resolution coverage and eliminating duplicate signal counting in the LLM taste-profiling layer.

AI/ML Engineering Intern, Dreamable Inc., CA

May 2025 to Aug 2025

- Trained, tuned, and deployed Machine Learning models in a production Cloud environment using Python, PyTorch, TensorFlow, and Hugging Face, working collaboratively with a cross-functional team.
- Fine-tuned the Qwen-2.5 7B LLM on GCP Cloud using LoRA to lower training cost while delivering a Q&A model with comparable accuracy within budget, tracking experiments with Weights & Biases.
- Curated NLP datasets using Python, Pandas, NumPy, and the Hugging Face Datasets library, cleaning and preprocessing training data to improve data quality for model fine-tuning.
- Developed a Generative AI outreach agent using LangChain and the OpenAI API to automate messaging workflows, cutting manual outreach time and increasing response rates.