

Kunj P. Shah

San Francisco, CA | +1 (628)-529-6990 | kunjcr2@gmail.com | [linkedin.com/kunjcr2](https://www.linkedin.com/in/kunjcr2) | github.com/kunjcr2 | kunjshah.vercel.app

F-1 international student eligible for CPT and OPT; eligible for H-1B change of status. The \$100,000 H-1B fee DOES NOT APPLY.

EDUCATION

San Jose State University

Jan 2026 - May 2027

B.S., Computer Science

San Jose

- **GPA:** 3.94/4.00
- **Achievements:** Dean's List, Researching on Next Sentence Prediction, CodePath Advanced DSA Batch
- **Coursework:** Object Oriented Programming, Advanced Data Structures and Algorithms, Computer Architecture

SKILLS

- **Programming:** Python, SQL, Java, Javascript, Git
- **Machine Learning & NLP:** Supervised & Unsupervised Learning, PyTorch, Tensorflow, Scikit-learn, Azure ML, Weights and Biases, Pandas, NumPy, Neural Network, Deep Learning, Reinforcement Learning, Langchain, n8n, Artificial Intelligence, Data Analysis
- **LLM:** Transformers, Large Language Models, Fine Tuning, Model Training, Model Inference, Model Alignment, Model Tracking and Experimenting, Model Evaluation, vLLM
- **Backend System:** Restful API, FastAPI, PostgreSQL, ORMs, Docker, Microsoft Azure Ecosystem, CI/CD, Agile, Microsoft Office Suite

EXPERIENCE

A10 Networking

May 2026 - Aug 2026

Deep Learning Intern - LLM Research and Model safety

San Jose, CA

- Researching Vision Language Model (VLM) guardrails to develop AI-driven safety mechanisms for multimodal systems on the AI Security team.

Routes Technologies

Oct 2025 - Jan 2026

AI Engineer Intern

Remote

- Architected an LLM-powered NL-to-SQL RAG pipeline with few-shot prompting, input normalization (470+ synonym mappings), parameterized SQL sanitization, session memory, and Weights & Biases observability — deployed across 5 Azure ML managed online endpoints.
- Built an SVD collaborative filtering recommendation pipeline on Azure ML (scikit-surprise), evaluated with NDCG@5 and Precision/Recall@K, paired with an LLM-driven taste profiler that batch-generates user preference vectors via a nightly ETL scheduled job.
- Engineered multi-source recipe ingestion: Scrapy web crawler with JSON-LD extraction, Instagram Graph API and TikTok oEmbed integrations with GPT-4o-mini recipe classification, lingua language filtering, and faster-whisper audio transcription fallback.

Dreamable Inc.

May 2025 - Aug 2025

AI/ML Engineering Intern

San Francisco, CA

- Fine-tuned the Qwen-2.5 7B model using PyTorch, TensorFlow, and Hugging Face on GCP Cloud Run, applying LoRA to lower training cost and memory usage while delivering a Q&A model with comparable accuracy within budget
- Curated NLP datasets using pandas, NumPy, and the Hugging Face Datasets library, cleaning and preprocessing training data to improve data quality and readiness for model fine-tuning
- Evaluated model hyperparameters through systematic experimentation, tracking evaluation loss and metrics with Weights & Biases to optimize model performance
- Developed an AI outreach agent using n8n, LangChain, Exa.ai, and the OpenAI API to automate messaging workflows, cutting manual outreach time, increasing response rates, and documenting processes in Microsoft Office Suite

PROJECTS

StableLM-2-1.6B End-to-end | [Github](#)

Jan 2026 - Feb 2026

Personal Project

San Francisco

- Built an end-to-end LLM post-pretraining pipeline using PyTorch, Transformers, and PEFT to fine-tune StableLM 1.6B on 140K UltraChat conversations via LoRA (r=256), with custom data preprocessing to convert multi-turn dialogues into single-turn SFT format
- Implemented GRPO alignment using TRL and a DeBERTa-v3 reward model to safety-align the fine-tuned LLM on 4K PKU-SafeRLHF samples, merging LoRA adapters and publishing final weights to HuggingFace for public inference
- Developed a production inference API using FastAPI and vLLM to serve the aligned model with configurable sampling parameters (temperature, top-k, top-p), containerized with Docker on a GPU-enabled vLLM base image for deployment

Kanting | [Github](#)

Oct 2025

CalHacks 12.0

San Francisco

- Built a Video RAG system using Python, Flask, and FAISS enabling natural language search over YouTube videos by extracting transcripts with Whisper (GPU), embedding them with Sentence Transformers (all-MiniLM-L6-v2, 384-dim), and indexing with FAISS for semantic retrieval
- Designed an end-to-end AI pipeline that retrieves relevant transcript segments and generates context-aware answers with GPT-4o-mini, returning timestamped video links for precise source attribution
- Developed a Flask-based API (video ingestion, querying, health, stats) with modular components for video downloading, vector storage, and LLM orchestration, runnable end-to-end on Google Colab with ngrok