

Kunj P. Shah

F-1 international student eligible for CPT and OPT; eligible for H-1B change of status. The \$100,000 H-1B fee DOES NOT APPLY.
San Francisco, CA | +1 6285296990 | kunjcr2@gmail.com | linkedin.com/kunjcr2 | github.com/kunjcr2 | kunjcr2.github.io

EDUCATION

San Jose State University

B.S., Computer Science

Jan 2026 - May 2027

San Jose

- **GPA:** 3.88/4.00
- **Achievements:** Dean's List, CodePath Advanced DSA Batch
- **Coursework:** Object Oriented Programing, Advanced Data Structures and Algorithms, Computer Architecture, Information Security, Machine Learning, Software Engineering

EXPERIENCE

A10 Networks

May 2026 - Aug 2026

Deep Learning Engineer Intern - AI Firewall Team

San Jose, CA

- Architected and trained TrooperViz, a multimodal safety guardrail model built with Python, PyTorch, and Hugging Face Transformers, enabling comprehensive evaluation of text, image-text inputs, multi-turn conversations, and assistant-generated outputs across diverse safety categories
- Fine-tuned a vision-language model for multimodal safety classification and merged it with a text-only guardrail model using SLERP, creating a unified safety system with broader coverage across visual, conversational, and generative-AI interactions.
- Engineered end-to-end data and evaluation workflows - including annotation, Pandas-based preprocessing, synthetic multi-turn conversation generation, model benchmarking, and interactive inference via Gradio and vLLM - streamlining the testing pipeline and accelerating safety assessments

Routes Technologies

Oct 2025 - Present

AI Engineer Intern

Remote

- Architected an LLM-powered NL-to-SQL RAG pipeline with few-shot prompting, input normalization (470+ synonym mappings), parameterized SQL sanitization, session memory, and Weights & Biases observability - deployed across 5 Azure ML managed online endpoints.
- Built an SVD collaborative-filtering recommendation pipeline on Azure ML, evaluated with precision, recall, and MAP metrics, and paired it with an LLM-driven taste profiler that batch-generates user preference vectors in a nightly ETL job, resulting in more relevant recommendations and faster cold-start handling
- Engineered a multi-source recipe ingestion system using a Scrapy crawler with JSON-LD extraction, Instagram Graph API and TikTok oEmbed integrations, GPT-4o-mini for recipe classification, lingua for language filtering, and faster-whisper for audio transcription fallback, which expanded the recipe catalog and cut ingestion latency

Dreamable Inc.

May 2025 - Aug 2025

AI/ML Engineering Intern

San Francisco, CA

- Fine-tuned the Qwen-2.5 7B model using PyTorch, TensorFlow, and Hugging Face on GCP Cloud Run, applying LoRA to lower training cost and memory usage while delivering a Q&A model with comparable accuracy within budget
- Curated NLP datasets using pandas, NumPy, and the Hugging Face Datasets library, cleaning and preprocessing training data to improve data quality and readiness for model fine-tuning
- Evaluated model hyperparameters through systematic experimentation, tracking evaluation loss and metrics with Weights & Biases to optimize model performance
- Developed an AI outreach agent using n8n, LangChain, Exa.ai, and the OpenAI API to automate messaging workflows, cutting manual outreach time, increasing response rates, and documenting processes in Microsoft Office Suite

SKILLS

- **Programming:** Python, SQL, Java, Javascript, Git
- **Machine Learning & NLP:** Supervised & Unsupervised Learning, PyTorch, Tensorflow, Scikit-learn, Azure ML, Weights and Biases, Pandas, NumPy, Neural Networks, Deep Learning, Reinforcement Learning, Langchain, n8n, Artificial Intelligence, Data Analysis
- **LLM:** Transformers, LLMs, Fine Tuning, Training, Inference, Alignment, Tracking and Experimenting, Evaluation
- **Backend System:** Restful API, FastAPI, PostgreSQL, ORMs, Docker, Microsoft Azure Ecosystem, CI/CD, Agile