Kunj P. Shah

F-1 international student eligible for CPT and OPT; eligible for H-1B change of status. The $100,000 H-1B fee DOES NOT APPLY.
San Francisco, CA | +1 6285296990 | kunjcr2@gmail.com | linkedin.com/kunjcr2 | github.com/kunjcr2 | kunjcr2.github.io

EDUCATION

San Jose State University **Jan 2026 - May 2027**
B.S., Computer Science *San Jose*

- **GPA:** 3.88/4.00
- **Achievements:** Dean's List, Researching on Next Sentence Prediction, CodePath Advanced DSA Batch
- **Coursework:** Object Oriented Programing, Advanced Data Structures and Algorithms, Computer Architecture, Information Security, Machine Learning, Software Engineering

SKILLS

- **Programming**: Python, SQL, Java, Javascript, Git
- **Machine Learning & NLP**: Supervised & Unsupervised Learning, PyTorch, Tensorflow, Scikit-learn, Azure ML, Weights and Biases, Pandas, NumPy, Neural Network, Deep Learning, Reinforcement Learning, Langchain, n8n, Artificial Intelligence, Data Analysis
- **LLM**: Transformers, Large Language Models, Fine Tuning, Model Training, Model Inference, Model Alignment, Model Tracking and Experimenting, Model Evaluation, vLLM
- **Backend System**: Restful API, FastAPI, PostgreSQL, ORMs, Docker, Microsoft Azure Ecosystem, CI/CD, Agile

EXPERIENCE

A10 Networks **May 2026 - August 2026**
Deep Learning Engineer Intern - AI Firewall Team San Jose, CA

- Developed a multimodal input/output guardrail system for VLM-powered applications, fine-tuning a compact language model with SFT and chain-of-thought rationales and serving production inference on Hopper GPUs via vLLM.
- Built a custom data annotation and verification pipeline to overcome the scarcity and noise of publicly available multimodal safety datasets, producing a curated training corpus with inter-annotator agreement metrics and automated quality filtering.
- Designed a systematic evaluation framework with safety policy diagnostics, collaborating within the AI Firewall team and using Weights & Biases for experiment tracking.

Routes Technologies **Oct 2025 - Jan 2026**
AI Engineer Intern *Remote*

- Architected an LLM-powered NL-to-SQL RAG pipeline with few-shot prompting, input normalization (470+ synonym mappings), parameterized SQL sanitization, session memory, and Weights & Biases observability - deployed across 5 Azure ML managed online endpoints.
- Built an SVD collaborative filtering recommendation pipeline on Azure ML (scikit-surprise), evaluated with NDCG@5 and Precision/Recall@K, paired with an LLM-driven taste profiler that batch-generates user preference vectors via a nightly ETL scheduled job.
- Engineered multi-source recipe ingestion: Scrapy web crawler with JSON-LD extraction, Instagram Graph API and TikTok oEmbed integrations with GPT-4o-mini recipe classification, lingua language filtering, and faster-whisper audio transcription fallback.

Dreamable Inc. **May 2025 - Aug 2025**
AI/ML Engineering Intern *San Francisco, CA*

- Fine-tuned the Qwen-2.5 7B model using PyTorch, TensorFlow, and Hugging Face on GCP Cloud Run, applying LoRA to lower training cost and memory usage while delivering a Q&A model with comparable accuracy within budget
- Curated NLP datasets using pandas, NumPy, and the Hugging Face Datasets library, cleaning and preprocessing training data to improve data quality and readiness for model fine-tuning
- Evaluated model hyperparameters through systematic experimentation, tracking evaluation loss and metrics with Weights & Biases to optimize model performance
- Developed an AI outreach agent using n8n, LangChain, Exa.ai, and the OpenAI API to automate messaging workflows, cutting manual outreach time, increasing response rates, and documenting processes in Microsoft Office Suite

PROJECTS

StableLM-2-1.6B End-to-end Github **Jan 2026 - Feb 2026**
Personal Project *San Francisco*

- Built an end-to-end LLM post-pretraining pipeline using PyTorch, Transformers, and PEFT to fine-tune StableLM 1.6B on 140K UltraChat conversations via LoRA (r=256), with custom data preprocessing to convert multi-turn dialogues into single-turn SFT format
- Implemented GRPO alignment using TRL and a DeBERTa-v3 reward model to safety-align the fine-tuned LLM on 4K PKU-SafeRLHF samples, merging LoRA adapters and publishing final weights to HuggingFace for public inference
- Developed a production inference API using FastAPI and vLLM to serve the aligned model with configurable sampling parameters (temperature, top-k, top-p), containerized with Docker on a GPU-enabled vLLM base image for deployment

MORE PROJECTS ON GITHUB